

CPSC 532W Homework 3

Naomi Graham

March 9, 2021

All the code can be found on: https://github.com/n6graham/cpsc532_hw3.

1 Importance sampling

1.1 Code

For Importance Sampling I used the evaluation-based implementation based off of Jason's HW2 code.

```
1 def evaluate_program(ast):
2     """Evaluate a program as desugared by daphne, generate a sample from the
3     prior
4     Args:
5         ast: json FOPPL program
6     Returns: sample from the prior of ast
7     """
8     PROCS = {} #program procedures
9     for i in range(len(ast)-1):
10         proc = ast[i]
11         proc_name, proc_arg_names, proc_expr = proc[1], proc[2], proc[3]
12         PROCS[proc_name] = (proc_arg_names, proc_expr)
13
14     #print(PROCS)
15     # expr is ast[-1]
16
17 def eval(expr, sigma, scope):
18     if is_const(expr, scope):
19         if type(expr) in [int, float]:
20             expr = torch.Tensor([expr]).squeeze()
21         return expr, sigma
22     elif is_var(expr, scope):
23         return scope[expr], sigma
24     elif is_let(expr, scope):
25         var_name, sub_expr, final_expr = expr[1][0], expr[1][1], expr[2]
26         var_value, sigma = eval(sub_expr, sigma, scope)
27         return eval(final_expr, sigma, {**scope, var_name: var_value})
28     elif is_if(expr, scope):
29         cond_expr, true_expr, false_expr = expr[1], expr[2], expr[3]
30         cond_value, sigma = eval(cond_expr, sigma, scope)
31         if cond_value:
32             return eval(true_expr, sigma, scope)
33         else:
34             return eval(false_expr, sigma, scope)
```

```

35     elif is_sample(expr, scope):
36         dist_expr = expr[1]
37         dist_obj, sigma = eval(dist_expr, sigma, scope)
38         return dist_obj.sample(), sigma
39     elif is_observe(expr, scope):
40         # need to do something special here
41         dist_expr, obs_expr = expr[1], expr[2]
42         dist_obj, sigma = eval(dist_expr, sigma, scope)
43         obs_value, sigma = eval(obs_expr, sigma, scope)
44         sigma['logW'] = sigma['logW'] + dist_obj.log_prob(obs_value)
45         return obs_value, sigma
46     else:
47         proc_name = expr[0]
48         consts = []
49         for i in range(1, len(expr)):
50             const, sigma = eval(expr[i], sigma, scope)
51             consts.append(const)
52         if proc_name in PROCS:
53             proc_arg_names, proc_expr = PROCS[proc_name]
54             new_scope = {**scope}
55             for i, name in enumerate(proc_arg_names):
56                 new_scope[name] = consts[i]
57             return eval(proc_expr, sigma, new_scope)
58         else:
59             return PRIMITIVES[proc_name](*consts), sigma
60
61
62     return eval(ast[-1], {'logW': 0}, {})
63

```

I also defined some extra functions: likelihood-weighting, compute-expectation, and compute-variance.

```

1  def likelihood_weighting(L, ast):
2      weighted_samples = []
3
4      for i in range(0, L):
5          r, sigma = evaluate_program(ast)
6          #r, sigma = eval(e, {'logW': 0}, [])
7          logW = sigma['logW']
8          weighted_samples.append((r, logW))
9
10     return weighted_samples
11

```

```

1  def compute_expectation(weighted_samples):
2      """
3      using a stream of weighted samples, compute according
4      to eq 4.6
5      """
6      L = len(weighted_samples)
7      log_weights = [weighted_samples[i][1] for i in range(0, L)]
8      weights = np.exp(np.array(log_weights))
9      print(weights)
10     r = [weighted_samples[i][0] for i in range(0, L)]
11     denom = sum(weights)
12     print("denominator is", denom)
13     numerator = sum([r[i] * weights[i] for i in range(0, L)])

```

```

14     #numerator = sum( [ weighted_samples[i][0] * weighted_samples[i][1] for i
    in range(0,L) ])
15
16     return numerator/denom
17
18
19 def compute_variance( weighted_samples , mu):
20
21     L = len( weighted_samples )
22     log_weights = [ weighted_samples[i][1] for i in range(0,L) ]
23     weights = np.exp(np.array(log_weights))
24     r = [ weighted_samples[i][0] for i in range(0,L) ]
25     denom = sum(weights)
26     numerator = sum( [ (torch.square(r[i]) - torch.square(mu)) * weights[i]
    for i in range(0,L) ])
27     return numerator/denom
28
29
30

```

1.2 results

The values returned are:

Program 1:

expectation is: tensor(7.2086) variance is tensor(0.7626)

Program 2

expectation is: tensor([2.1571, -0.5566]) variance is tensor([0.0561, 0.8358])

Program 3:

expectation is: tensor(0.7429) variance is tensor(0.1910)

Program 4:

expectation is: tensor(0.3221) variance is tensor(0.2183)

2 MH within Gibbs

2.1 code

For MH withing Gibbs I used a graph-based implementation based off of Jason's HW2 code.

```

1 def deterministic_eval(exp):
2     "Evaluation function for the deterministic target language of the graph
    based representation."
3     if type(exp) is list:
4         op = exp[0]
5         args = exp[1:]
6         return env[op]( *map(deterministic_eval , args))
7     elif type(exp) in [int , float]:
8         # We use torch for all numerical objects in our evaluator
9         return torch.Tensor([ float(exp) ]).squeeze()
10    elif type(exp) is torch.Tensor:
11        return exp
12    elif type(exp) is bool:
13        return torch.tensor(exp)
14    else:
15        print("expression is:", exp)
16        print(type(exp))
17        raise Exception("Expression type unknown.", exp)

```

```

18
19 def topological_sort(nodes, edges):
20     result = []
21     visited = {}
22     def helper(node):
23         if node not in visited:
24             visited[node] = True
25             if node in edges:
26                 for child in edges[node]:
27                     helper(child)
28             result.append(node)
29     for node in nodes:
30         helper(node)
31     return result[::-1]
32
33 def plugin_parent_values(expr, trace):
34     if type(expr) == str and expr in trace:
35         return trace[expr]
36     elif type(expr) == list:
37         return [plugin_parent_values(child_expr, trace) for child_expr in expr]
38     else:
39         return expr
40
41 def sample_from_joint(graph):
42     """This function does ancestral sampling starting from the prior."""
43     # TODO insert your code here
44     """
45     1. Run topological sort on V using V and A, resulting in an array of v's
46     2. Iterate through sample sites of the sorted array, and save sampled
47        results on trace dictionary using P and Y
48        - If keyword is sample*, first recursively replace sample site names with
49          trace values in the expression from P. Then, run deterministic_eval.
50        - If keyword is observe*, put the observation value in the trace
51        dictionary
52     3. Filter the trace dictionary for things sample sites you should return
53     """
54     procs, model, expr = graph[0], graph[1], graph[2]
55     nodes, edges, links, obs = model['V'], model['A'], model['P'], model['Y']
56     sorted_nodes = topological_sort(nodes, edges)
57
58     sigma = {}
59     trace = {}
60     for node in sorted_nodes:
61         keyword = links[node][0]
62         if keyword == "sample*":
63             link_expr = links[node][1]
64             link_expr = plugin_parent_values(link_expr, trace)
65             dist_obj = deterministic_eval(link_expr)
66             trace[node] = dist_obj.sample()
67         elif keyword == "observe*":
68             trace[node] = obs[node]
69
70     expr = plugin_parent_values(expr, trace)
71     return deterministic_eval(expr), sigma, trace

```

I added the function MH-Gibbs which also contains functions accept and Gibbs-step.

```

1  def MH_Gibbs(graph , numsamples):
2      model = graph[1]
3      vertices = model['V']
4      arcs = model['A']
5      links = model['P'] # link functions aka P
6
7      # sort vertices for ancestral sampling
8      V_sorted = topological_sort(vertices , arcs)
9
10     def accept(x, cX, cXnew, Q):
11         # compute acceptance ratio to decide whether
12         # we keep cX or accept a new sample/trace cXnew
13         # cX and cXnew are the proposal mappings (dictionaries)
14         # which assign values to latent variables
15
16         # cXnew corresponds to the values for the new samples
17
18         # take the proposal distribution for the current vertex
19         # this is Q(x)
20         Qx = Q[x][1]
21
22         # we will sample from this with respect to cX and cXnew
23
24         # the difference comes from how we evaluate parents
25         # plugging into eval
26         p = plugin_parent_values(Qx,cX)
27         pnew = plugin_parent_values(Qx,cXnew)
28
29         # p = Q(x)[X := \mathcal{X}]
30         # p' = Q(x)[X := \mathcal{X}']
31         # note that in this case we only need to worry about
32         # the parents of x to sample from the proposal
33
34
35         # evaluate
36         d = deterministic_eval(p) # d = EVAL(p)
37         dnew = deterministic_eval(pnew) #d' = EVAL(p')
38
39         ### compute acceptance ratio ###
40
41         # initialize log alpha
42         logAlpha = dnew.log_prob(cXnew[x]) - d.log_prob(cX[x])
43
44         ###  $V_x = \{x\} \cup \{v: x \in PA(v)\}$  ###
45         startindex = V_sorted.index(x)
46         Vx = V_sorted[startindex:]
47
48         # compute alpha
49         for v in Vx:
50             Pv = links[v] # getting a bug here
51             v_exp = plugin_parent_values(Pv,cX) #same as we did for p and
52             v_exp_new = plugin_parent_values(Pv,cXnew)
53             dv_new = deterministic_eval(v_exp_new[1])
54             dv = deterministic_eval(v_exp[1])
55
56             ## change below
57

```

```

58         logAlpha = logAlpha + dv_new.log_prob(cXnew[v])
59         logAlpha = logAlpha - dv.log_prob(cX[v])
60     return torch.exp(logAlpha)
61
62
63 def Gibbs_step(cX,Q):
64     # here we need a list of the latent (unobserved) variables
65     Xobsv = list(filter(lambda v: links[v][0] == "sample*", V_sorted))
66
67     for u in Xobsv:
68         # here we are doing the step
69         # d <- EVAL(Q(u) [X := \cX])
70         # note it suffices to consider only the non-observed variables
71         Qu = Q[u][1]
72         u_exp = plugin_parent_values(Qu,cX)
73         dist_u = deterministic_eval(u_exp).sample()
74         cXnew = {**cX}
75         cX[u] = dist_u
76
77         #compute acceptance ratio
78         alpha = accept(u,cX,cXnew,Q)
79         val = Uniform(0,1).sample()
80
81         if val < alpha:
82             cX = cXnew
83     return cX
84
85
86 Q = links # initialize the proposal with P (i.e. using the prior)
87 cX_list = [ sample_from_joint(graph)[2] ] # initialize the state/trace
88
89 for i in range(1,numsamples):
90     cX_0 = {**cX_list[i-1]} #make a copy of the trace
91     cX = Gibbs_step(cX_0,Q)
92     cX_list.append(cX)
93
94     samples = list(map(lambda cX: deterministic_eval(plugin_parent_values(
95         graph[2], cX)), cX_list))
96     #samples = [ deterministic_eval(plugin_parent_values(graph[2],X)) for X
97     in cX_list ]
98
99     return samples

```

2.2 Results

Here are the results from running MH Gibbs for programs 1 through 4:

```
==== running program 1 ====
Program 1 mean: 6.4470267
Program 1 variance: 0.026307987
Total run time: 26.21826195716858

==== running program 2 ====
Program 2 slope mean is: 1.6526821
Program 2 bias mean is: 1.3684516
Program 2 slope variance is: 0.07386868
Program 2 bias variance is: 0.61285996
Total run time: 55.65074110031128

==== running program 3 ====
Program 3 mean is: 0.9986666666666667
[Program 3 variance is: 0.001331555555555556
total run time: 579.4399008750916

==== running program 4 ====
Program 4 mean is: 2.5e-05
Program 4 variance is: 2.4999375e-05
total run time: 575.306489944458_
```

2.2.1 Program 1




Figure 1: trace plot (left) and histogram (right)

2.2.2 Program 2




Figure 2: Program 2 (bias): trace plot (left) and histogram (right)

2.2.3 Program 3

2.2.4 Program 4




Figure 3: Program 2 (slope): trace plot (left) and histogram (right)




Figure 4: Program 3 histogram




Figure 5: Program 4 histogram